



Random forests

Jeffrey Leek, Assistant Professor of Biostatistics
Johns Hopkins Bloomberg School of Public Health

Random forests

similar to bagging

1. Bootstrap samples
2. At each split, bootstrap variables
3. Grow multiple trees and vote
or average

at each potential split, only a subset of variables is considered.

Pros:

1. Accuracy


Cons:

1. Speed `slow`
2. Interpretability
3. Overfitting
so it's very important to use cross validation
when using random forests

Random forests


each tree is based on a separate subset of the data

(randomly
sampled)



The ensemble model

$$\text{Forest output probability } p(c|\mathbf{v}) = \frac{1}{T} \sum_t p_t(c|\mathbf{v})$$



<http://www.robots.ox.ac.uk/~az/lectures/ml/lect5.pdf>

this link doesn't work

Iris data

```
data(iris); library(ggplot2)
inTrain <- createDataPartition(y=iris$Species,
                                p=0.7, list=FALSE)
training <- iris[inTrain,]
testing <- iris[-inTrain,]
```

Random forests

```
library(caret)
modFit <- train(Species~ .,data=training,method="rf",prox=TRUE)
modFit                                     random forest
```

```
105 samples
  4 predictors
  3 classes: 'setosa', 'versicolor', 'virginica'
```

No pre-processing

Resampling: Bootstrap (25 reps)

Summary of sample sizes: 105, 105, 105, 105, 105, 105, ...

Resampling results across tuning parameters:

mtry	Accuracy	Kappa	Accuracy SD	Kappa SD
2	0.9	0.9	0.03	0.04
3	0.9	0.9	0.03	0.05
4	0.9	0.9	0.03	0.05

Getting a single tree

```
getTree(modFit$finalModel,k=2)  get the second tree
```

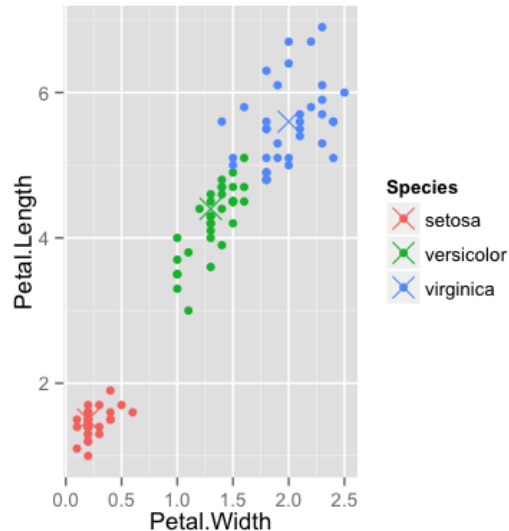
which variable
was used and
to split where

left daughter	right daughter	split var	split point	status	prediction	
1	2	3	4	0.70	1	0
2	0	0	0	0.00	-1	1
3	4	5	4	1.70	1	0
4	6	7	3	4.95	1	0
5	8	9	3	4.85	1	0
6	0	0	0	0.00	-1	2
7	10	11	4	1.55	1	0
8	12	13	1	5.95	1	0
9	0	0	0	0.00	-1	3
10	0	0	0	0.00	-1	3
11	0	0	0	0.00	-1	2
12	0	0	0	0.00	-1	2
13	0	0	0	0.00	-1	3

each row is 1 split

Class "centers"

```
irisP <- classCenter(training[,c(3,4)], training$Species, modFit$finalModel$prox)
irisP <- as.data.frame(irisP); irisP$Species <- rownames(irisP)
p <- qplot(Petal.Width, Petal.Length, col=Species, data=training)
p + geom_point(aes(x=Petal.Width, y=Petal.Length, col=Species), size=5, shape=4, data=irisP)
```



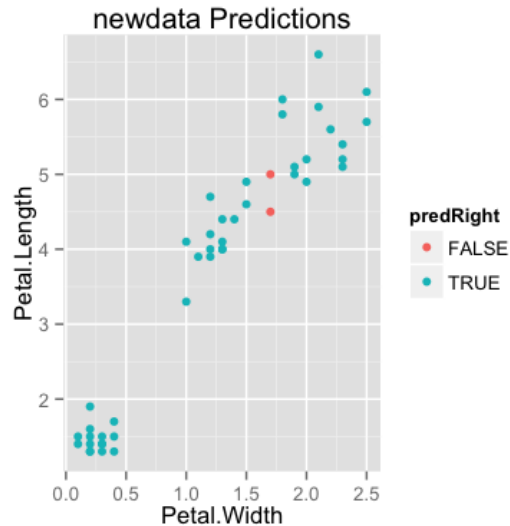
Predicting new values

```
pred <- predict(modFit,testing); testing$predRight <- pred==testing$Species  
table(pred,testing$Species)
```

pred	setosa	versicolor	virginica
setosa	15	0	0
versicolor	0	14	1
virginica	0	1	14

Predicting new values

```
qplot(Petal.Width,Petal.Length,colour=predRight,data=testing,main="newdata Predictions")
```



Notes and further resources

Notes:

- Random forests are usually one of the two top performing algorithms along with boosting in prediction contests.
- Random forests are difficult to interpret but often very accurate.
- Care should be taken to avoid overfitting (see [rfcv](#) function)

Further resources:

- [Random forests](#)
- [Random forest Wikipedia](#)
- [Elements of Statistical Learning](#)